

ID5030 Assignment 1

Haran Rajkumar

ED15B020

1. Linear Regression with gradient descent

1. General Code

```
from matplotlib import pyplot as plt
import random as rd
import pandas as pd
from functions import *

#Generates a sample hypothesis
def sample_hypothesis1(data_size,noise=0):
    x = np.arange(0,2*np.pi,2*np.pi/data_size)
    y = np.sin(x)
    if noise!=0:
        print("noise added")
        mean=np.std(y)
        for i in rd.sample(range(0,len(x)),int(len(x)*noise)):
            # y[i]+=mean*3*(i%2-0.5)
            y[i]=np.cos(x[i])
    return x,y

#Weight Initialisation
def initialize_weights(number_of_features,val=10**-5):
    w=np.full(number_of_features,val)
    return w

#Extrapolates Polynomial Features from the given 2D function
def create_features(x,number_of_features):
    X = np.zeros((len(x),number_of_features))
    for j in range(0,number_of_features):
        for i,a in enumerate(x):
            X[i][j]=a**j
    return X

#Separate Dataset into training and test datasets
```

```

def separate_datasets(X,y,train_percent=0.8):
    train_indices =
rd.sample(range(0,len(X)),int(train_percent*len(X)))
    test_indices = [i for i in range(0,len(X)) if i not in
train_indices]
    x_training = np.ones((len(train_indices),X.shape[1]))
    x_test = np.ones((len(test_indices),X.shape[1]))
    y_test = np.ones(len(test_indices))
    y_training = np.ones(len(train_indices))
    # print(len(train_indices),X.shape)
    for idx,i in enumerate(train_indices):
        x_training[idx,:]=X[i,:]
        y_training[idx]=y[i]

    for idx,i in enumerate(test_indices):
        x_test[idx,:]=X[i,:]
        y_test[idx]=y[i]

    return x_training,y_training,x_test,y_test

def batch_grad(X_training,y_training,alpha,w,grad_threshold):
    number_of_features=len(w)
    data_size=len(y_training)
    grad_w=np.full(number_of_features,11)
    L=np.array([0,1])
    count = 1
    while (abs(L[count]-L[count-1])>grad_threshold):
        Y_training=np.dot(X_training,w.T)
        grad_w=np.dot(-1*2*(y_training-Y_training).T,X_training)
        w=w-alpha*grad_w
        count+=1

val=sum((y_training-np.dot(X_training,w.T))**2)/float(data_size)
    L=np.append(L,val)
    return w

def calc_error(X_test,y_test,w):
    return sum((y_test-np.dot(X_test,w.T))**2)/len(y_test)


#Main function
alpha=10**-6
number_of_features=4
data_size = 10000
grad_threshold=10**-4

```

```
x,y= sample_hypothesis1(data_size,0)
X= create_features(x,number_of_features)
X_training,y_training,X_test,y_test= separate_datasets(X,y)

w = initialize_weights(number_of_features)
w=batch_grad(X_training,y_training,alpha,w,grad_threshold)
mean_squared_error=calc_error(X_test,y_test,w)
print("Error: %f"%(mean_squared_error))
```

2. I used a 2D - curve as the sample hypothesis (function 'sample_hypothesis'). I extrapolated features by taking the polynomial values of the samples (function 'create_features()'). Therefore, I am trying to fit the 2D curve with an nth order polynomial with n being the number of features. I plotted the 2D curve along with the predicted curve to see if it fits properly. I used the 'mean squared error'(function calc_error') as the test metric. It's a general test metric because it is normalised to the number of observations.



3. I used a criterion which computes the mean squared error and stops if the difference in errors of the current iteration to the previous one is lesser than a threshold (typically around 10^{-3} - 10^{-6}). I tried various learning rates and found that above a certain a value the prediction blows up and below a certain value it takes a very long time to converge. For the above hypothesis, a learning rate between $[10^{-5}, 10^{-6}]$ gave reasonable results.

2. Stochastic gradient descent and variants

I used the same test I used for testing Batch Gradient Descent, i.e. check if the predicted curve fits the graph and check if the error is below a reasonable limit. I've provided the graphs which compares the actual function to the predicted function

a. Stochastic gradient descent

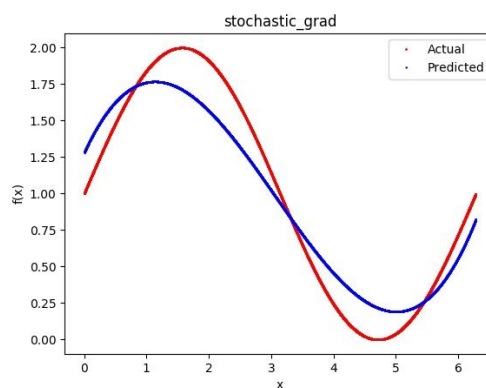
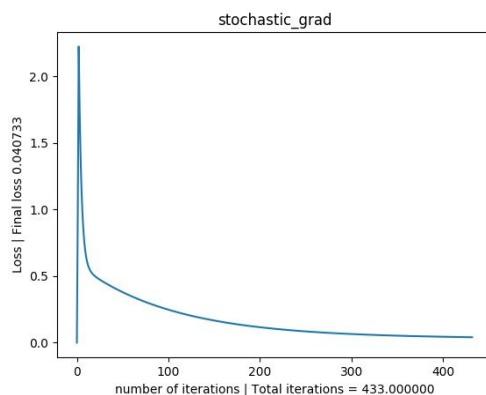
Formulae

Code

```
def stochastic_grad(X_training,y_training,alpha,w,grad_threshold):
    number_of_features=len(w)
    data_size=len(y_training)
    grad_w=np.full(number_of_features,1)
    L=np.array([0,1])
    count = 1
    while (abs(L[count]-L[count-1])>grad_threshold):
        for idx in range(0,data_size):
            Y_training=np.dot(X_training[idx,:],w.T)
            grad_w =
np.dot(-1*2*(y_training[idx]-Y_training).T,X_training[idx,:])
            w=w-alpha*grad_w
            count+=1

    val=sum((y_training-np.dot(X_training,w.T))**2)/float(data_size)
    L=np.append(L,val)
    return w
```

Plot



Merits

- It's more accurate compared to Batch gradient descent
-

Demerits

- Slower compared to the algorithms due to the static learning rate

b. Stochastic gradient descent with momentum

Formulae

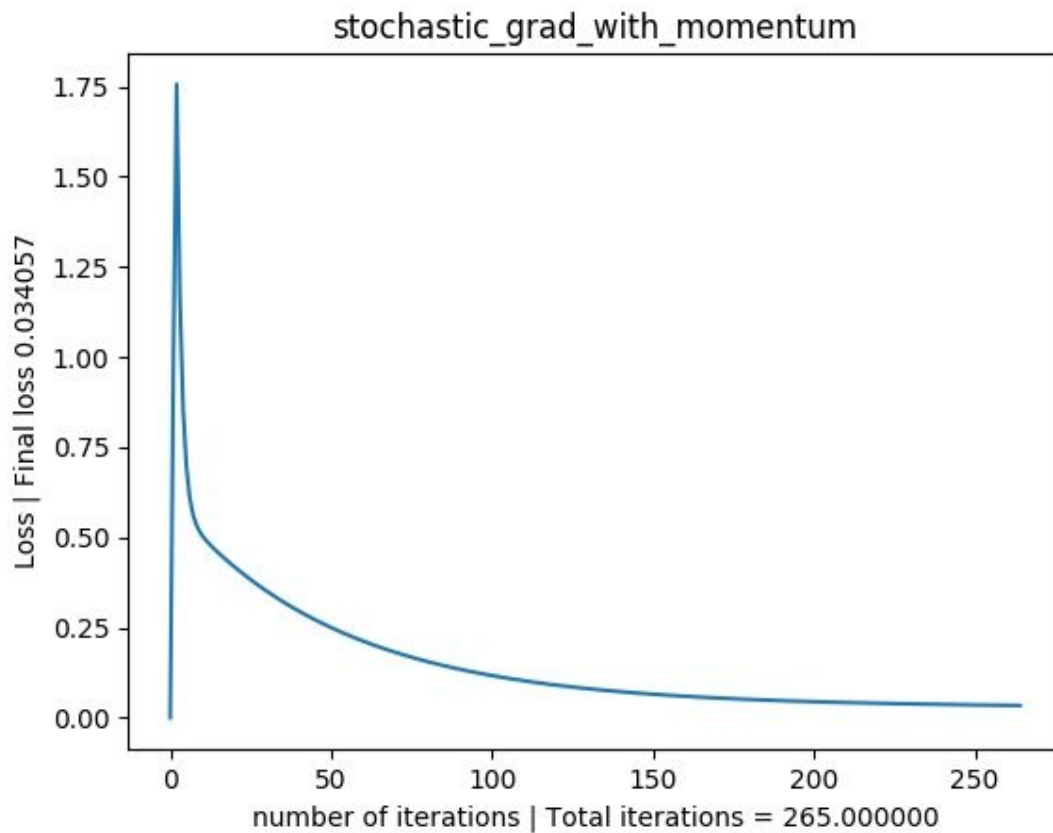
Code

```
def
stochastic_grad_with_momentum(X_training,y_training,alpha,w,grad_thresho
ld,eta=1):
    number_of_features=len(w)
    data_size=len(y_training)
    grad_w_dash=0
    grad_w=np.full(number_of_features,11)

    L=np.array([0,1])
    count = 1
    while (abs(L[count]-L[count-1])>grad_threshold):
        for idx in range(0,data_size):
            Y_training=np.dot(X_training[idx,:],w.T)
            grad_w =
np.dot(-1*2*(y_training[idx]-Y_training).T,X_training[idx,:])
            delta_w=-alpha*grad_w
            w=w + delta_w*eta - alpha*grad_w
            count+=1

    val=sum((y_training-np.dot(X_training,w.T))**2)/float(data_size)
    L=np.append(L,val)
    return w
```

Plots



Merits

-

Demerits

-

c. Adagrad

Formulae

Code

```
def adagrad(X_training,y_training,alpha,w,grad_threshold,eta=1):
    number_of_features=len(w)
    data_size=len(y_training)
    ada_g=np.zeros(number_of_features)
    grad_w=np.zeros(number_of_features)

    L=np.array([0,1])
    count = 1
    while (abs(L[count]-L[count-1])>grad_threshold):
        for idx in range(0,data_size):
```

```

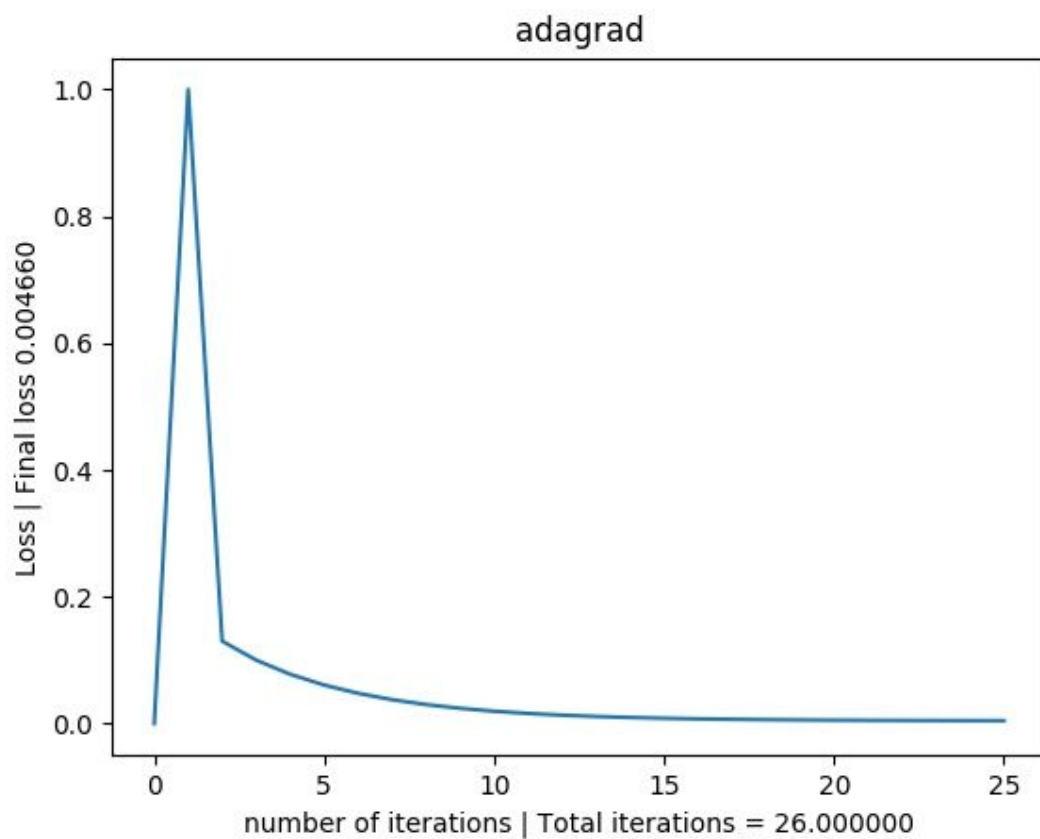
        Y_training=np.dot(X_training[idx,:],w.T)
        grad_w =
np.dot(-1*2*(y_training[idx]-Y_training).T,X_training[idx,:])
        ada_g+=grad_w**2
        w = w - eta*grad_w/np.sqrt(ada_g)

    count+=1

val=sum((y_training-np.dot(X_training,w.T))**2)/float(data_size)
    L=np.append(L,val)
    return w

```

Plots



Merits

-

Demerits

-

d. RMSprop

Formulae

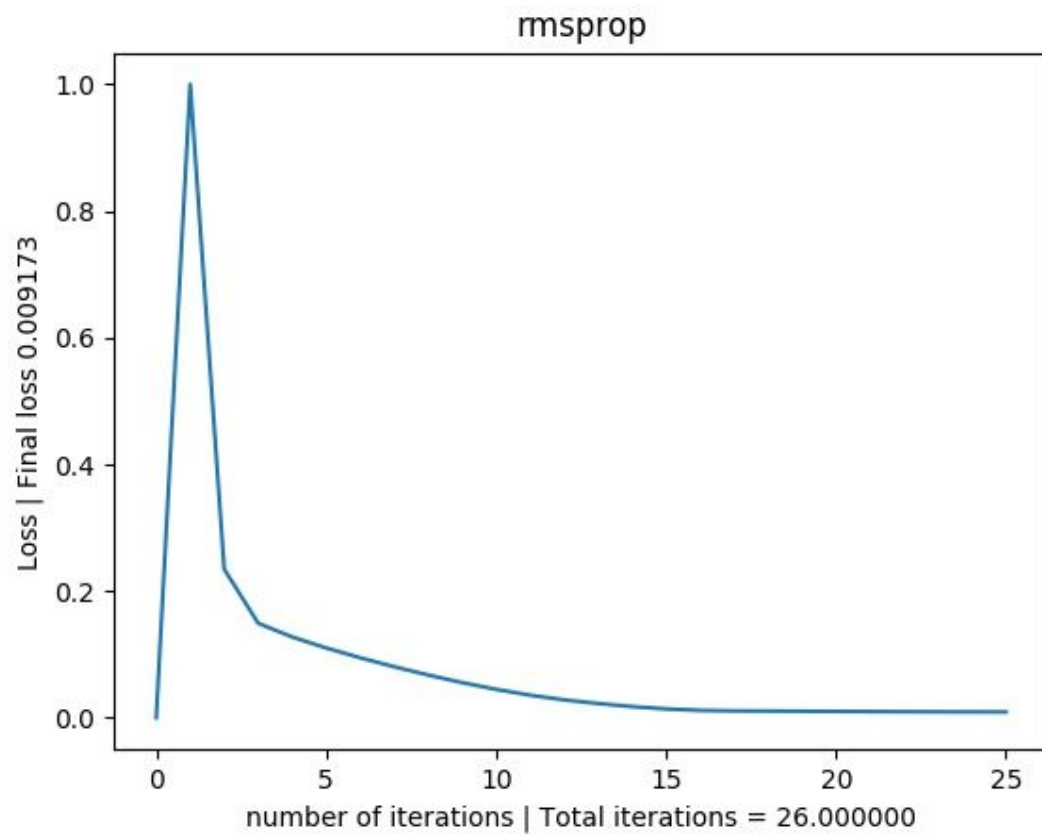
Code

```
def
rmsprop(X_training,y_training,alpha,w,grad_threshold,eta=10**-3,gamma=0.
9):
    number_of_features=len(w)
    data_size=len(y_training)
    v=np.zeros(number_of_features)
    grad_w=np.zeros(number_of_features)

    L=np.array([0,1])
    count = 1
    while (abs(L[count]-L[count-1])>grad_threshold):
        for idx in range(0,data_size):
            Y_training=np.dot(X_training[idx,:],w.T)
            grad_w =
np.dot(-1*2*(y_training[idx]-Y_training).T,X_training[idx,:])
            v=gamma*v+(1-gamma)*(grad_w**2)
            w = w - eta*grad_w/np.sqrt(v)
        count+=1

    val=sum((y_training-np.dot(X_training,w.T))**2)/float(data_size)
    L=np.append(L,val)
    return w
```


Plots



Merits

-

Demerits

-